Assignment 01

Learning from Data, Related Challenges, Linear Models for Regression

submitted for

EN3150 - Pattern Recognition

Department of Electronic and Telecommunication Engineering
University of Moratuwa

Udugamasooriya P. H. J.
220658U

12 August 2025

1 Impact of Outliers on Linear Regression

Question 02 We represent the independent variables in a matrix

$$\mathbf{X} = \begin{pmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_n \end{pmatrix},$$

the dependent variables in a vector

$$\mathbf{y} = \begin{pmatrix} y_1 & \cdots & y_n \end{pmatrix}^\top,$$

and directly use the result that

$$\mathbf{w}_{\text{OLS}} = \begin{pmatrix} w_0 & w_1 \end{pmatrix}^\top = \arg\min_{\mathbf{w}} (\mathbf{y} - \mathbf{X}\mathbf{w})^2 = \left(\mathbf{X}^\top \mathbf{X}\right)^{-1} \mathbf{X}^\top \mathbf{y}.$$

This is exactly what we do in code X, and the results obtained from it are as follows:

```
Ordinary Least Squares Weights (w): [ 3.91672727 -3.55727273]
```

Hence,

$$\mathbf{w}_{\text{OLS}} = \begin{pmatrix} 3.91672727 \\ -3.55727273 \end{pmatrix},$$

and the predicted linear model is

$$y = 3.91672727 - 3.55727273x.$$

A plot of the given data points against the predicted values is shown in Figure 1.

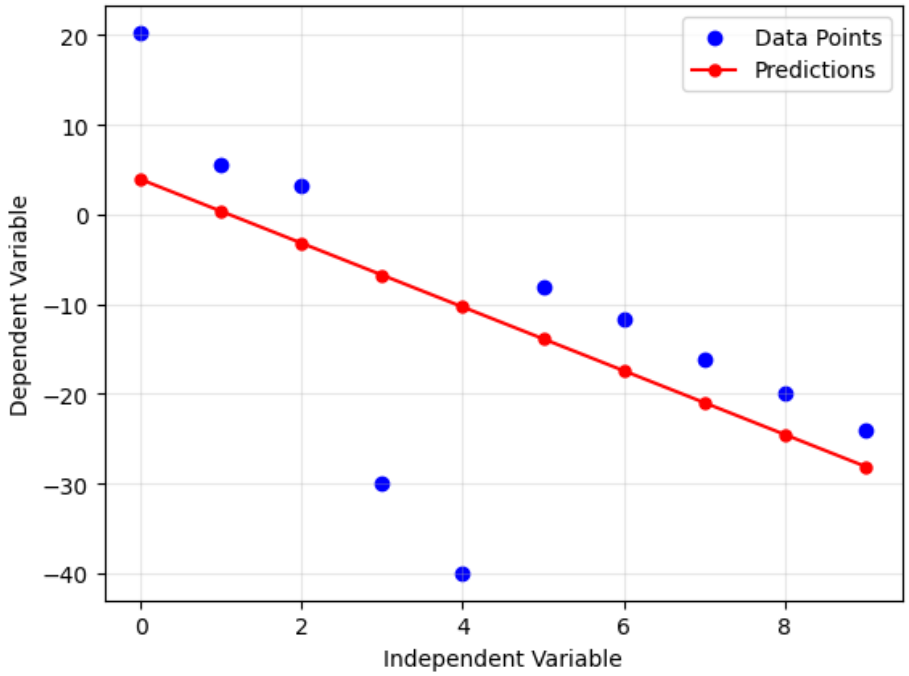

Figure 1: abc

Question 04 The code in X was used to calculate the loss for each model for each given value of β. The output of the code was the following:

```
Model 1 : [12 -4]
  Loss for beta = 1      : 0.435416262490386
  Loss for beta = 1e-06  : 0.9999999998258207
  Loss for beta = 1000.0 : 0.0002268287498440988
Model 2 : [ 3.91 -3.55]
  Loss for beta = 1      : 0.9728470518681676
  Loss for beta = 1e-06  : 0.9999999999999718
  Loss for beta = 1000.0 : 0.00018824684654645654
```

Summarizing these results in a table, we have

β	Model 1 Loss	Model 2 Loss
1	0.4354	0.9728
10^{-6}	0.9999	1.0000
10^{3}	0.0002	0.0002

Table 1: s

Question 05 We propose setting $\beta = 1$ to mitigate the impact of outliers.

With very small values of β, the squared error term starts to dominate, and the loss becomes approximately equal to 1, i.e.,

$$\frac{(y_i - \hat{y}_i)^2}{(y_i - \hat{y}_i)^2 + \beta^2} \approx \frac{(y_i - \hat{y}_i)^2}{(y_i - \hat{y}_i)^2} = 1,$$

making the result almost independent of the model used, and making it difficult to distinguish between several models.

Very large values of β would cause the loss to be approximately proportional to the squared error and very close to 0, i.e.,

$$\frac{(y_i - \hat{y_i})^2}{(y_i - \hat{y_i})^2 + \beta^2} \approx \frac{(y_i - \hat{y_i})^2}{\beta^2} \approx 0,$$

again making it difficult to distinguish between models. Further, inimizing the loss in this case would yield approximately the same result as that of minimizing the mean squared error.

It can be seen from the results above that $\beta = 10^3$ is too large, as the resulting losses from both models are both approximately equal, and very small and close to zero, making it difficult to distinguish between the two models.

We can also see that $\beta = 10^{-6}$ is too small, as the resulting losses from the models in this case are again both approximately equal but this time close to 1, leading to the same problem as above.

Hence, $\beta = 1$ is the best choice of the given options, as it has yielded comparable losses for both models.

Question 06 We will fix $\beta = 1$. The loss for Model 1 then, is 0.4354, whereas the loss for Model 2 is 0.9728. Clearly, Model 1 has a lower loss and is therefore the better/more suitable model.

Question 07 Let us start by rewriting the loss for a single data point as follows:

$$L\left(y_i, \hat{y_i}, \theta, \beta\right) = \begin{cases} \dfrac{1}{1 + \frac{\beta^2}{(y_i - \hat{y_i})^2}}, & y_i \neq \hat{y_i}, \\ 0, & \text{otherwise.} \end{cases}$$

It is clear then that for all $\hat{y_i}$, the loss is always non-negative and less than 1, with small values of $(y_i - \hat{y_i})^2$ resulting in losses close to 0 and larger values resulting in losses close to 1; and importantly, this property also holds true for outliers.

This has the effect of keeping the loss restricted to a finite range, especially in the presence of outliers, in which case a simpler loss such as the squared error would become very large, and not be bounded above.

Further, because the loss is always non-negative, the mean loss over all the data points is minimized when the loss from each individual data point is as small as possible.

Suppose we choose some threshold $E \in (0, 1)$ such that we would like to have

$$\frac{1}{1 + \frac{\beta^2}{(y_i - \hat{y_i})^2}} \leq E.$$

This is equivalent to requiring

$$(y_i - \hat{y_i})^2 \leq \frac{E}{1 - E} \cdot \beta^2, \quad \text{or} \quad |y_i - \hat{y_i}| \leq \sqrt{\frac{E}{1 - E}} \cdot \beta.$$

Note that this defines an interval of values centered around the actual value y_i, that the predicted $\hat{y_i}$ might lie within, for which the loss can still be considered "small enough", and β allows us to control how big or small this interval is. β can be chosen big enough to include the large distance that the outliers are away from what one might expect their true values to be.

Hence, with this modification in place, we prevent the outliers from introducing very large losses, and encourage the model to focus more on minimzing the loss due to the inliers, which would contribute more significantly to the mean loss, due to the larger number of inliers compared to outliers.

Question 08 different loss

2 Loss Functions

Question 01 We calculate the squared error

$$\text{SE}(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

and binary cross entropy

$$\text{BCE}(\hat{y}_i, y_i) = -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i)$$

of each predicted value $\hat{y}_i$ against the given corresponding target value y_i.

True Value (y_i)	Predicted Value ($\hat{y_i}$)	SE($\hat{y}_i, y_i$)	BCE($\hat{y}_i, y_i$)
1	0.005	0.9900	5.2983
1	0.010	0.9801	4.6052
1	0.050	0.9025	2.9957
1	0.100	0.8100	2.3026
1	0.200	0.6400	1.6094
1	0.300	0.4900	1.2040
1	0.400	0.3600	0.9163
1	0.500	0.2500	0.6931
1	0.600	0.1600	0.5108
1	0.700	0.0900	0.3567
1	0.800	0.0400	0.2231
1	0.900	0.0100	0.1054
1	1.000	0.0000	0.0000
	Mean	0.4402	1.4407

Table 2: Losses for each predicted value

A plot of the different losses against the predicted values is shown in Figure 2.

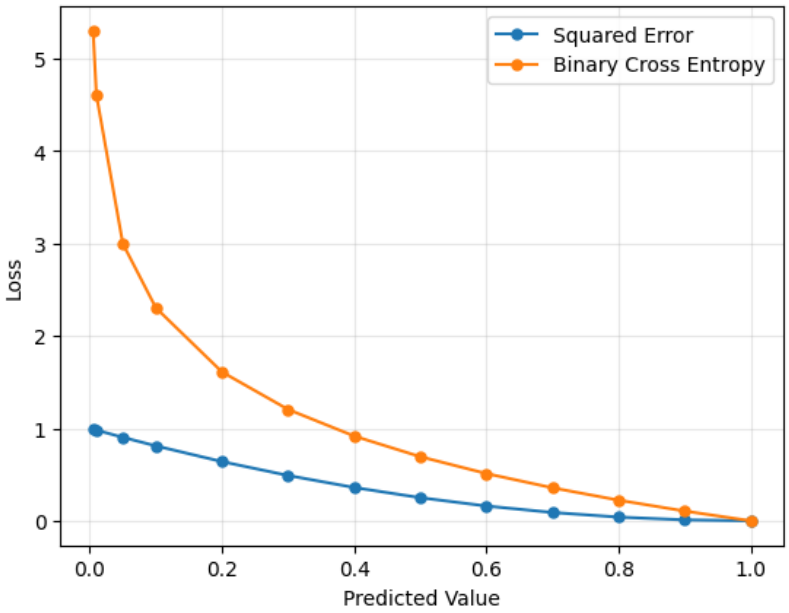

Figure 2: abc

Question 02 The MSE loss is more appropriate for Application 1 and the BCE loss is more appropriate for Application 2.

For Application 1, where the quantity being predicted is a continuous value, we would like the loss to vary gradually with difference between the predicted value and true value.

For Application 2, where the goal is to map a continuous value to a binary value based on a certain threshold, we would like continuous values below the threshold to have a significantly larger loss than values above the threshold (or vice versa). In this case, we are not very interested in how far away the predicted continuous value is away from the threshold.

3 Data Pre-Processing

Question 01 We start by visualizing the distribution of each feature value to decide on a suitable method of scaling. Figure 3 contains plots obtained of the distribution of each feature in the dataset.

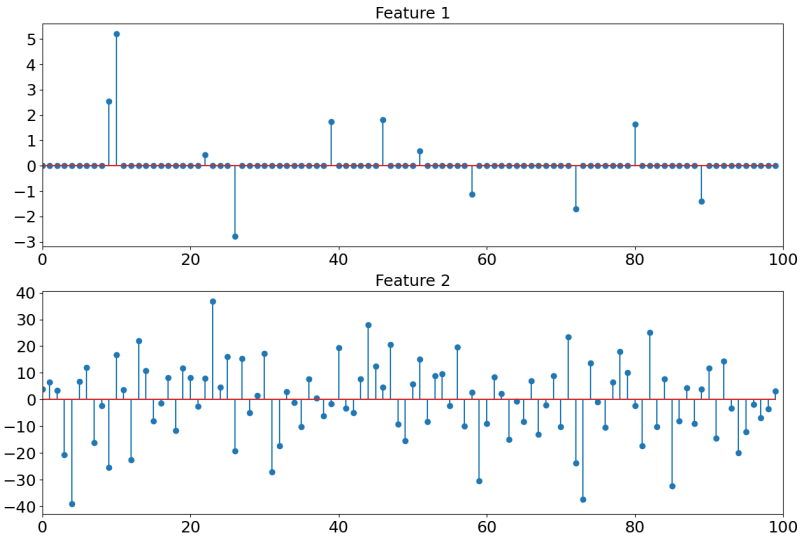

Figure 3: abc

We then run the code in Y to obtain the following summary statistics of each feature:

```
Feature 1
  Mean     : 0.06963158220374253
  Standard Deviation : 0.751690461643816
  Maximum    : 5.2
  Minimum    : -2.790493210023752
  Range     : 7.990493210023752
Feature 2
  Mean     : -0.45935709567298505
  Standard Deviation : 14.351150951654933
  Maximum    : 36.752574877667975
  Minimum    : -39.11938330852965
  Range     : 75.87195818619762
```

Based on the plots and the summary statistics above, we conclude the following;

- both features have means close to zero
- the features vary on different scales, as they have significantly different standard deviations
- both features take on both positive and negative values
- Feature 1 is sparsely distributed, with most values being equal to 0

To bring the values of both features to a similar scale while still preserving the structure and properties of each feature, we consider the three following scaling methods;

1. standard scaling,
2. min-max scaling, and
3. max-abs scaling.

We rule out min-max scaling as it would limit both feature values to a range between 0 and 1, affecting the “symmetric” variation of both feature values among both negative and positive values.

To choose between standard and max-abs scaling, we consider the sparsity of Feature 1.

We observe that standard scaling would map zeros of Feature 1 to non-zero values, resulting in a loss of sparsity—a propertsy one would likely want to preserve. On the other hand, max-abs scaling would not affect sparsity, as it would map zeros to zeros. It also maps to a range between -1 and 1, so negative values map to negative values, and positives to positives.

Hence, we choose max-abs scaling as the method of scaling for both feature values. A plot of both feature values after the above scaling was applied is shown in Figure **??**. The recalculated summary statistics for the scaled features are as follows:

```
vab
```